

Llama roughly doubles answer length; GPT-OSS stays close to the original

